

**STRATEGY & OPERATIONS**

Vytalix Decision Framework (Build / Test / Defer / Kill)

Vytalix company-build documentation · reference
20_DECISION_FRAMEWORK

Status labels apply throughout: KNOWN, VERIFIED, ASSUMPTION, ESTIMATE, TO VALIDATE, PROPOSED. Nothing in this document claims traction, approvals, partners or revenue that do not exist. Governing file: FACTS_BASE.

PREPARED FOR

The Founder, Vytalix

DATE

8 September 2026

VERSION

1.0 · Confidential draft

Contents

Executive view	1
1. The ten criteria	2
2. Thresholds and verdicts	3
3. The eight live ideas — scoring table	4
4. Scoring rationale by idea	5
5. Portfolio view and what the scores say about the next 90 days	6
6. Blank template for future ideas	7
Priorities / Risks / Next actions	8

Owner: Founder (CEO), with the Chief of Staff as scorer of record · **Status:** PROPOSED · **Version:** 1.0 (8 September 2026) · **Governed by:** /FACTS_BASE.md (every score is an ASSUMPTION or ESTIMATE unless the evidence column says KNOWN or VERIFIED) · **Used by:** /docs/18_90_DAY_EXECUTION_PLAN.md gates G-C (Day 3), G-D (Day 60), task 3.16 (grant bid/no-bid, Fri 23 Oct 2026), task 4.7 (MVP build decision, Tue 8 Dec 2026) and task 4.16 (quarterly re-score); /docs/16_ORGANISATION_AND_SOPS.md SOP 03 step 3 and SOP 07; /docs/17_FOUNDER_OPERATING_SYSTEM.md §2.4 (quarterly re-score) · **Related:** /docs/00_MASTER_AUDIT.md §5-§6, /docs/01_BUSINESS_MODEL.md , /docs/04_MATERNALINK_PRODUCT_STRATEGY.md , /docs/05_MARKET_STRATEGY.md , /assets/PRICING_STRUCTURE.md

Executive view

- One scoring sheet for every idea that competes for the founder's time or the company's cash: ten criteria, each 1–10, unweighted, total out of 100, with three criteria reverse-scored (cost, regulatory risk, execution complexity — a high score means low cost, low risk, low complexity) and four knock-out rules that override the total. Verdicts: **BUILD ≥ 70** · **TEST 50–69** · **DEFER 35–49** · **KILL < 35**.
- Applied to the eight live ideas today: **BUILD** founder-led advisory (76). **TEST** the research report as lead magnet (65), MaternaLink v1 care-coordination (59) and UK NHS pilot pursuit as a design-partner evaluation (54). **DEFER** e-learning launch in year 1 (49) and white-label MaternaLink for NGOs (45). **KILL** the MaternaLink AI risk score as a product now (33; knock-out on regulatory risk) and the Ekiti proposal at £300 per woman per year as written (33; knock-out on strategic fit and an unvalidated price). Every score is a desk ESTIMATE from the audit, not a forecast.
- "KILL" means the idea as stated is not pursued; it does not mean the underlying opportunity is dead. The AI risk score re-enters as a v3 research track with a data partner and an ethics route; the Ekiti conversation re-enters as a partner-aware, re-priced option scored at gate G-D on Day 60 (Thu 12 Nov 2026).
- "TEST" is a commitment to a defined experiment with a budget cap, a KPI and a decision date — not a soft yes. The three TEST verdicts already have their experiments in the 90-day plan (report 1 by Week 9–10; PRD sign-off and the MVP build decision on Day 86; 15 UK discovery interviews and five design-partner briefs by Day 90).
- Re-score quarterly (first: week of Mon 14 December 2026) and at every gate; any score that moves by three or more points must cite new evidence. A blank template (§6) is provided for future ideas; scoring without evidence labels is not accepted.

1. The ten criteria

#	Criterion	What it measures	Direction	1 means	10 means	Evidence that raises confidence
1	Strategic fit	Alignment with the group thesis: services fund the company, MaternaLink (in development) is the equity story, UK-first then Nigeria, non-diagnostic v1, truth-first	Normal	Contradicts the thesis or the language rules	Directly advances the thesis and at least two pillars	Business model §1; FACTS_BASE A4–A6
2	Revenue potential	Realistic cash in 12–24 months and the size of the stream at maturity, probability-weighted	Normal	No path to cash	Large, recurring, high-margin, evidenced	Signed engagements; validated pricing; comparable programmes
3	Customer demand	Evidence that a named buyer wants this and can pay	Normal	No buyer conversation; no public evidence	Multiple paying customers or signed LOIs	Discovery calls; LOIs; public procurement signals

#	Criterion	What it measures	Direction	1 means	10 means	Evidence that raises confidence
4	Speed to market	Time to first paying customer or first live use	Normal	> 24 months	< 1 month	Dependency chain in the 90-day plan
5	Cost	Cash required to reach the first result (reverse-scored)	Reverse	> £250k before first result	< £2k	Budget lines in the 90-day plan and financial model
6	Scalability	Growth without proportional founder time or cost	Normal	Founder-hours only	Software-like margins across markets	Unit economics in the business model
7	Competitive advantage	Defensible difference versus named alternatives, buildable within 12 months	Normal	Commodity; incumbents stronger on every axis	Unique, protected, evidenced	Competitive intelligence doc
8	Regulatory risk	Exposure to medical-device, data-protection, clinical-safety, procurement and financial-promotion rules (reverse-scored)	Reverse	Requires a claim or pathway Vytalix cannot lawfully make or fund now	No regulatory exposure beyond ordinary B2B	Legal and regulatory doc \$2–\$6
9	Execution complexity	Number of dependencies, partners, skills and unknowns between now and the first result (reverse-scored)	Reverse	Many external dependencies, several unknowns, no owner	One owner, known steps, no external dependency	Task dependencies in the 90-day plan
10	Founder leverage	Whether the solo founder can deliver it with current skills, time and network without hiring first	Normal	Requires a team the company does not have	Founder can deliver alone within the time-allocation rule	Founder OS \$1; hiring plan triggers

Reverse-scoring rule. Criteria 5, 8 and 9 are assessed as High / Medium / Low first, then scored so that Low = 8–10, Medium = 4–7, High = 1–3. This keeps every criterion pointing the same way: a higher total is always better, and a reader can compare ideas by a single number.

Evidence labels. Each score carries one of KNOWN / VERIFIED / ASSUMPTION / ESTIMATE / TO VALIDATE. A score of 8 or above on criteria 2, 3 or 7 requires at least one KNOWN or VERIFIED item in the evidence column; without it the score is capped at 7.

2. Thresholds and verdicts

Verdict	Total /100	Meaning	What happens next	Rationale for the band
BUILD	≥ 70	Commit founder time and budget; put it in the operating plan with owner, dates and KPIs	Enters the 90-day plan as P0/P1 tasks; reported weekly	An average of 7 across all ten criteria means the idea is strong on nearly every dimension and weak on none that matters; with unweighted criteria this is the level at which a pre-seed company can commit scarce time without hedging
TEST	50–69	Promising but with real unknowns; run a bounded experiment	A written test card: hypothesis, budget cap (£ ESTIMATE), KPI, decision date; re-scored on that date	The mid-band is where most product ideas sit before validation; the right response is to buy evidence cheaply, not to commit or abandon
DEFER	35–49	Not now; revisit when a named trigger fires	Logged with the trigger (e.g., "after MVP", "after first NHS design partner", "Month 9"); no time spent beyond a waitlist or a note	Below the midpoint the idea fails on several criteria at once; spending time now would starve a BUILD or TEST item

Verdict	Total /100	Meaning	What happens next	Rationale for the band
KILL	< 35	Not pursued as stated	Recorded with the reason; the underlying opportunity may re-enter as a differently stated idea	An average under 3.5 means the idea fails on most dimensions; carrying it forward costs credibility with investors and buyers

Knock-out rules (override the total):

- **K1 — Strategic fit $\leq 3 \rightarrow$ KILL** regardless of total. An idea that contradicts the thesis or the language rules is not improved by scoring well elsewhere.
- **K2 — Regulatory risk $\leq 2 \rightarrow$ cannot be BUILD or TEST.** Maximum verdict DEFER until a lawful, funded pathway exists; if the idea requires a claim prohibited by `/FACTS_BASE.md` §4 it is KILL as stated.
- **K3 — Founder leverage ≤ 2 with no funded team $\rightarrow$ cap at DEFER.** The company has one operator; an idea that needs a team the company cannot pay for waits for the round or the revenue.
- **K4 — Customer demand $\leq 2 \rightarrow$ cannot be BUILD.** Nothing is built without at least one buyer conversation that says so.

Tie-breaks and adjustments: if two ideas compete for the same founder hours, the one that creates revenue sooner (criterion 4 $\times$ criterion 2) wins; if a score sits on a band boundary, the verdict takes the lower band until evidence moves it.

3. The eight live ideas — scoring table

All scores are ESTIMATES made on 8 September 2026 from the audit and the strategy documents. Columns 1–10 follow §1; 5, 8 and 9 are reverse-scored (high = good).

Idea	1 Fit	2 Rev	3 Demand	4 Speed	5 Cost (R)	6 Scale	7 Adv	8 Reg (R)	9 Complex (R)	10 Leverage	Total	Knock-out	Verdict
A. Founder-led advisory	9	7	6	10	10	3	5	9	9	8	76	—	BUILD
B. Vytalix research report as lead magnet	8	3	6	8	8	5	6	8	6	7	65	—	TEST
C. MaternaLink v1 care-coordination (non-diagnostic)	10	7	5	4	4	8	5	7	4	5	59	—	TEST
D. UK NHS pilot pursuit (design-partner evaluation first)	9	5	6	3	6	8	4	5	3	5	54	—	TEST
E. E-learning launch in year 1	6	4	3	4	6	9	2	7	5	3	49	—	DEFER
F. White-label MaternaLink for NGOs	6	5	5	2	3	8	4	5	3	4	45	—	DEFER
G. MaternaLink AI risk score as a product now	6	3	4	1	2	7	6	1	1	2	33	K2, K3	KILL (now)

Idea	1 Fit	2 Rev	3 Demand	4 Speed	5 Cost (R)	6 Scale	7 Adv	8 Reg (R)	9 Complex (R)	10 Leverage	Total	Knock-out	Verdict
H. Ekiti proposal at £300/woman/year as written	4	5	2	2	5	4	2	4	2	3	33	K4 (and K1 borderline)	KILL (as written)

4. Scoring rationale by idea

Each row: score · reasoning · evidence label.

A. Founder-led advisory (Advisory days A1, retainers A2, packages C1–C10 where proof exists) — 76, BUILD

Criterion	Score	Reasoning	Evidence
Strategic fit	9	The cash story in FACTS_BASE A6; funds the product; builds references and market knowledge	KNOWN (strategy)
Revenue potential	7	£15k–£30k contracted by Day 90 (ESTIMATE); day rates £1,200–£1,800; capped by founder days	ESTIMATE
Customer demand	6	S1/S2 segments have short cycles and a clear pain (NHS readiness, investor readiness); zero conversations held yet	ASSUMPTION; TO VALIDATE via 40 touches in Week 1
Speed to market	10	Sellable the week the entity, insurance and a contract template exist	KNOWN (dependency chain)
Cost (R)	10	Near-zero cash: templates, insurance, CRM free tier	KNOWN
Scalability	3	Founder-hours; associates add margin but not scale	KNOWN (business model)
Competitive advantage	5	Crowded market; the UK–Nigeria corridor and maternity focus are distinctive but unproven	ASSUMPTION
Regulatory risk (R)	9	Ordinary B2B consultancy; PI insurance; no clinical claims	KNOWN
Execution complexity (R)	9	One owner, known steps, SOP 01 written	KNOWN
Founder leverage	8	Founder can deliver alone; a technical or clinical associate needed only for C3/C5/C6/C8/C9	ASSUMPTION on founder expertise (TO CONFIRM domain depth)

Verdict — BUILD. This is the only idea that scores 8 or above on speed, cost, regulatory simplicity, execution and founder leverage at once, and it is the only one that produces cash inside 90 days. Its weakness — scalability at 3 — is accepted because its job is to fund the company and earn references, not to be the exit. The 90-day plan already treats it as P0 (sales quota from Day 8, first SOW Day 30, three by Day 90). Demand is the unvalidated score: if fewer than five discovery calls happen by Day 14 or no proposal converts by Day 45, the re-score reduces criterion 3 and the package mix is revisited, but the verdict holds as long as the day-rate floor of £1,200 is not breached.

B. Vytalix research report as lead magnet ("The State of Digital Maternal Care Coordination: UK and Nigeria, 2026") — 65, TEST

Criterion	Score	Reasoning	Evidence
Strategic fit	8	Credibility asset for every pillar; feeds proposals, investor materials and the PRD; master audit opportunity 6	KNOWN
Revenue potential	3	Indirect only; 5 inbound conversations targeted; no direct sales	ESTIMATE

Criterion	Score	Reasoning	Evidence
Customer demand	6	Buyers and investors read landscape reports; no evidence yet that this one will be downloaded	ASSUMPTION
Speed to market	8	Draft by Week 7, publish Week 9–10	KNOWN (plan)
Cost (R)	8	£1k–£2k design; founder time for 20 interviews	ESTIMATE
Scalability	5	One asset, reused many times; not a revenue engine	KNOWN
Competitive advantage	6	Few small-n mixed UK–Nigeria reports exist on this question; transparent method is the differentiator	ASSUMPTION
Regulatory risk (R)	8	Professional interviews, no patient data; consent, ICO registration and privacy notice needed	KNOWN
Execution complexity (R)	6	20 interviews, a survey, a landscape scan, design and PR — all founder-run	KNOWN
Founder leverage	7	Founder can deliver; a designer for £1k–£2k	ESTIMATE

Verdict — TEST. The report is the cheapest credibility asset available and it doubles as the market-validation work the company must do anyway (the 15 UK interviews are shared with task 3.7). It stays a TEST rather than a BUILD because its value is indirect and unproven: the test card is "150 downloads in 14 days, 400 in 90 days, 5 inbound conversations, 3 citations" with a decision on report 2 taken at the quarterly re-score in December. If downloads fall below 50 in 14 days, the second report is not commissioned and the interview material is reused as articles instead.

C. MaternaLink v1 care-coordination platform (non-diagnostic; care-coordination, communication, education, monitoring-support) — 59, TEST

Criterion	Score	Reasoning	Evidence
Strategic fit	10	The equity story, re-based to the intended-use statement adopted at gate G-C	KNOWN
Revenue potential	7	Site licences £40k–£120k/yr (ESTIMATE) in the UK; programme licences £8k–£60k in Africa; pilots small; 12–24 months away	ESTIMATE
Customer demand	5	National reviews (Ockenden, Kirkup) evidence the communication problem; no buyer has yet said it would pay; Ekiti outcome unknown	VERIFIED (problem) / TO VALIDATE (willingness to pay)
Speed to market	4	MVP 12–16 weeks after a funded build decision (Day 86 at the earliest); first live use 2027	KNOWN (plan)
Cost (R)	4	MVP build, CSO, DPIA, security baseline, hosting — a six-figure envelope over 12 months (ESTIMATE; see MVP spec)	ESTIMATE
Scalability	8	Software; multi-market; SMS channel for Nigeria	ASSUMPTION
Competitive advantage	5	Care-coordination and handover focus with interoperability is a defensible position, but incumbents exist in both markets (competitive intelligence doc)	ASSUMPTION
Regulatory risk (R)	7	Designed to sit outside device classification (TO VALIDATE with counsel); still needs DTAC, DCB0129, DSPT, DPIA, NDPA in Nigeria	KNOWN (pathway) / TO VALIDATE (qualification)
Execution complexity (R)	4	IP resolution, prototype audit, CSO, clinical advisers, design partners, engineering — many dependencies	KNOWN
Founder leverage	5	Needs a fractional CTO, a CSO and clinical advisers before build; founder can run discovery and PRD	KNOWN

Verdict — TEST. MaternaLink v1 is the reason the company exists and it scores 10 on fit, but it cannot be a BUILD today: IP is undocumented (gate G-B), the prototype is unaudited, no buyer has confirmed willingness to pay and no clinician

has reviewed the PRD. The test is already written into the plan — PRD v0.2 signed off with a CSO and a clinical adviser (Day 60), 15 UK discovery interviews and five design-partner briefs (Day 90), the prototype audit (Day 30) and the MVP build decision on Tue 8 Dec 2026 (build / build-lite / wait). Criterion 3 rises to 7 or more only with a signed design-partner LOI; criterion 9 rises once the IP deed is executed and the audit reports "pass" or "conditions". A total of 70 or more at the December re-score is the condition for a funded BUILD.

D. UK NHS pilot pursuit (design-partner evaluation with one trust or private maternity provider, then a paid pilot) — 54, TEST

Criterion	Score	Reasoning	Evidence
Strategic fit	9	UK-first thesis; a trust evaluation is the credibility milestone investors ask for	KNOWN
Revenue potential	5	Design partnership £0–£5k; paid pilot £20k–£60k after evaluation; site licence later (all ESTIMATE)	ESTIMATE
Customer demand	6	Problem evidenced by national reviews; trusts have been burnt by unvalidated maternity apps; no UK conversation held yet	VERIFIED (problem) / TO VALIDATE
Speed to market	3	Procurement and governance 8–16 weeks after artefacts exist; artefacts (DTAC pack, hazard log, DPIA, Cyber Essentials) are Phase 3–4 work	KNOWN
Cost (R)	6	Mostly founder time; CSO days; Cyber Essentials fee	ESTIMATE
Scalability	8	One successful evaluation is repeatable across trusts with the same artefacts	ASSUMPTION
Competitive advantage	4	Incumbent maternity digital tools exist; Vytalix's advantage is artefacts-first and coordination focus, unproven	ASSUMPTION
Regulatory risk (R)	5	DCB0160 on the trust side, DTAC, DSPT, information-governance sign-off; manageable but non-trivial	KNOWN
Execution complexity (R)	3	Depends on IP, MVP or a demonstrable prototype, CSO, a clinical champion and a trust innovation lead	KNOWN
Founder leverage	5	Founder can run discovery and relationship building; delivery needs CSO and engineering	KNOWN

Verdict — TEST. The right idea at the wrong altitude: a "pilot pursuit" today would be a promise without a product. Scored as what the plan actually does — discovery interviews, a synthesis note, five design-partner briefs and one LOI in drafting by Day 90 — it is a bounded test with near-zero cash cost. It stays below BUILD because speed (3) and complexity (3) are structural until the MVP decision and the assurance artefacts exist. Trigger to re-score upward: a written expression of interest from a maternity digital lead or a trust innovation team plus a CSO engaged. The rule from the outreach engine applies throughout: never describe any NHS conversation as a partnership or an approval.

E. E-learning launch in year 1 (courses on sale before Month 9) — 49, DEFER

Criterion	Score	Reasoning	Evidence
Strategic fit	6	A named pillar, but the business model places first revenue at Month 9 and treats it as year-2 income	KNOWN
Revenue potential	4	£49–£299 per course (ESTIMATE); volume requires an audience that does not exist yet	ESTIMATE
Customer demand	3	No waitlist, no learner conversation, no employer or NGO buyer identified	KNOWN (absence)
Speed to market	4	Two courses could be authored in a quarter, but the audience takes longer than the content	ESTIMATE
Cost (R)	6	LMS £25–£250/month, production time; modest cash, significant founder hours	ESTIMATE
Scalability	9	Self-serve, high gross margin once built	KNOWN (model)

Criterion	Score	Reasoning	Evidence
Competitive advantage	2	Crowded market of platforms and professional bodies; no accreditation may be claimed	KNOWN
Regulatory risk (R)	7	Low if no CPD or accreditation claims are made and clinical content is reviewed	KNOWN
Execution complexity (R)	5	Content authoring, clinical review, platform, payments, marketing funnel	KNOWN
Founder leverage	3	Every hour spent here is taken from the revenue block; founder cannot author, review and market alone in year 1	KNOWN

Verdict — DEFER. Scalability is the best in the set and the regulatory burden is light, but demand is unevidenced, the competitive position is weak and the founder's hours are the binding constraint. The plan's treatment stands: a waitlist page and two course outlines by Day 90 (task 4.14), free primers from Month 6 to build the funnel, first courses at Month 9 (ESTIMATE), corporate and NGO licences later. Trigger to re-score: 500 waitlist sign-ups or one corporate or NGO training conversation with a budget — either would move criteria 3 and 2 enough for TEST.

F. White-label MaternaLink for NGOs (branded patient-facing layer licensed to NGO programmes) — 45, DEFER

Criterion	Score	Reasoning	Evidence
Strategic fit	6	Consistent with the NGO route to a Nigeria cohort, but white-labelling a product that does not exist yet is a year-2 stream in the business model	KNOWN
Revenue potential	5	Set-up £25k–£75k plus £2–£8 per user per year (ESTIMATE); real but distant	ESTIMATE
Customer demand	5	NGOs have enrolled women, CHEWs and donor reporting needs; none has been asked whether they want a white-label rather than a programme licence	ASSUMPTION
Speed to market	2	Requires v1 to exist, plus multi-tenant branding and data-residency configuration	KNOWN
Cost (R)	3	Adds engineering on top of the MVP envelope	ESTIMATE
Scalability	8	Software licence; repeatable across programmes and telecoms later	ASSUMPTION
Competitive advantage	4	Incumbent programmes (mDoc, MomConnect-style services) already serve NGO and state cohorts	VERIFIED (incumbents exist)
Regulatory risk (R)	5	Cross-border data flows, NDPA registration, safeguarding on shared phones; manageable with design	KNOWN
Execution complexity (R)	3	Depends on v1, an NGO partner, in-country associate, data-residency architecture	KNOWN
Founder leverage	4	Founder can hold NGO conversations; delivery needs the product team	KNOWN

Verdict — DEFER. The idea is a delivery model for a product that is itself still a TEST, so it cannot score higher than the product it depends on. The nearer path to an NGO cohort is the programme licence in the product strategy §13 (enrolment, SMS workflows, defaulter tracing, logframe exports) priced to a grant budget line, which needs no white-label engineering. Trigger to re-score: an NGO or telecom partner asking for a branded instance in writing after v1 is live. Until then the white-label option is a line in the partnership deck, not a workstream.

G. MaternaLink AI risk score ("Maternal Instability Score") as a product now — 33, KILL (now); re-enters as the v3 research track

Criterion	Score	Reasoning	Evidence
Strategic fit	6	Long-term vision yes; as a product now it contradicts FACTS_BASE A4 and the language rules	KNOWN

Criterion	Score	Reasoning	Evidence
Revenue potential	3	No lawful sale before validation and a device pathway; years away	KNOWN
Customer demand	4	Clinicians want earlier warning; buyers will not procure an unvalidated predictive tool	ASSUMPTION
Speed to market	1	Data partnership, ethics approval, validation study, SaMD conformity assessment: > 24 months	KNOWN
Cost (R)	2	Research, data governance, clinical evaluation, Approved Body costs: well above £250k (ESTIMATE)	ESTIMATE
Scalability	7	If ever validated, high	ASSUMPTION
Competitive advantage	6	Multimodal (clinical + voice + social) concept is distinctive on paper; no data, model or validation exists	VERIFIED (concept only)
Regulatory risk (R)	1	Almost certainly a medical device (prediction/prognosis intended purpose); no MHRA pathway, ethics approval or data access exists	KNOWN
Execution complexity (R)	1	Requires a university or NHS research partner, data-sharing agreements, ethics, a clinical lead, ML engineering and QMS	KNOWN
Founder leverage	2	Cannot be delivered by the current team	KNOWN

Verdict — KILL as a product now (knock-outs K2 and K3). Shipping or even pitching a predictive score today would expose Vytalix to medical-device regulation it cannot meet, to investor diligence it would fail, and to reputational harm in a field where over-claiming is unforgivable. This is the "KILL-now" referenced at gate G-C. The opportunity is preserved, not discarded: the concept moves to the v3 research track in the product strategy §8, advancing only through a university or NHS research partnership with data access, ethics approval and a clinical lead, funded by research grants (NIHR i4i, Grand Challenges — eligibility TO VALIDATE). It is re-scored as a research proposal, not a product, when such a partner exists.

H. The Ekiti State proposal at £300 per woman per year as written (15,000 women, £4.5m Phase 1) — 33, KILL (as written); a re-scoped option is scored at gate G-D

Criterion	Score	Reasoning	Evidence
Strategic fit	4	Nigeria is the second market, but a £36m headline government sale led by an unresolved collaborator contradicts UK-first, truth-first and the IP position	KNOWN
Revenue potential	5	Large headline; probability-weighted value is low — no response, no budget line, unvalidated price	VERIFIED (proposal only)
Customer demand	2	No evidence of acceptance, contract or payment since 18 March 2026; the state already runs a funded programme	VERIFIED
Speed to market	2	Government cycles 3–9 months at best; this one has been silent for six months	KNOWN
Cost (R)	5	Pursuing it costs founder time and relationship capital more than cash; delivering it would cost far more than the company has	ESTIMATE
Scalability	4	Government programmes scale politically, not commercially; donor logic dominates	ASSUMPTION
Competitive advantage	2	mDoc's Digital Mom Project is incumbent, donor-funded, with ~24,000 women onboarded (public reporting)	VERIFIED
Regulatory risk (R)	4	NDPA, procurement and anti-bribery exposure; IP not owned; board approval required for any government proposal	KNOWN
Execution complexity (R)	2	Depends on IP resolution, in-country delivery capacity, a product that does not exist, and a ministry that has not replied	KNOWN
Founder leverage	3	The relationship, if real, sits with David Agunede; the founder cannot deliver a state programme alone	KNOWN

Verdict — KILL as written (knock-out K4; K1 borderline). The proposal as sent — £300 per woman per year, £4.5m for Phase 1, £36m for the state — is not restated under any option because the price is an unvalidated proposal figure, the state has an incumbent, and Vytalix does not yet own what was proposed. That is a decision about a document, not about Nigeria. The plan keeps the conversation alive on facts: an Ekiti status call by Day 26, an mDoc complementarity conversation by Day 54, a two-state scorecard (Oyo and Kano lead on the desk estimate in the market strategy \$2.1) by Day 55, and gate G-D on Thu 12 Nov 2026. If the ministry identifies a gap and mDoc is open to complementarity, a partner-aware LGA pilot at a £6–£40 per woman per year platform layer (ESTIMATE) and an envelope under £0.5m per year is drafted for board approval and scored on this sheet as a new idea; the default if that fails is to move state with an NGO or donor cohort.

5. Portfolio view and what the scores say about the next 90 days

Rank	Idea	Total	Verdict	Owner of the next step	Decision date
1	Founder-led advisory	76	BUILD	Founder	Continuous; first SOW Tue 13 Oct 2026
2	Research report as lead magnet	65	TEST	Founder (+ designer)	Report 2 go/no-go at the December re-score
3	MaternaLink v1 care-coordination	59	TEST	Founder + CSO + fractional CTO	MVP build decision Tue 8 Dec 2026
4	UK NHS pilot pursuit (design-partner first)	54	TEST	Founder	Design-partner pipeline review Fri 11 Dec 2026
5	E-learning launch in year 1	49	DEFER	Founder	Trigger: 500 waitlist sign-ups or one B2B training conversation with budget
6	White-label MaternaLink for NGOs	45	DEFER	Founder	Trigger: written request after v1 is live
7	AI risk score as a product now	33	KILL (now)	—	Re-enters only as a research proposal with a data partner
8	Ekiti proposal at £300/woman/year	33	KILL (as written)	Founder; board	Re-scoped option scored at G-D, Thu 12 Nov 2026

Reading: the portfolio is one cash engine (A), three bounded experiments that share most of their inputs (B, C, D all draw on the same 15–20 interviews, the same CSO and the same intended-use statement), two parked options with named triggers, and two ideas retired as stated so that the company never has to defend them in diligence. Founder time in the 90 days follows this ranking: revenue first (40%), then the shared validation work behind B–D.

Grant bid/no-bid (task 3.16, Fri 23 Oct 2026) uses the same sheet on each candidate call: a bid scores BUILD only if eligibility is KNOWN, a required clinical partner is at "agreed in principle" and the founder-day cost (10–15 days per bid, ESTIMATE) does not breach the sales quota for more than one week.

6. Blank template for future ideas

Copy this block into the decision log or into this file under a new heading. Every score needs a one-line reason and an evidence label; totals without reasons are not accepted at the Friday review.

IDEA: [name, stated precisely – "X for Y priced at Z by [date]"]
Proposed by: [name] Date scored: [dd Mon yyyy] Scorer: [name]
Trigger for scoring: [gate / inbound approach / quarterly re-score / other]
Depends on: [ideas or gates this cannot precede, e.g. G-B IP resolution, MVP build]

#	Criterion	Raw view (H/M/L for 5, 8, 9)	Score 1-10	One-line reason
1	Strategic fit			
2	Revenue potential			
3	Customer demand			
4	Speed to market			
5	Cost (reverse: low = high)			
6	Scalability			
7	Competitive advantage			
8	Regulatory risk (reverse)			
9	Execution complexity (reverse)			
10	Founder leverage			
	TOTAL /100			

Knock-outs checked: K1 fit ≤ 3 [] K2 reg ≤ 2 [] K3 leverage ≤ 2 without funded team [] K4 demand ≤ 2 []
Caps applied: 8+ on criteria 2/3/7 without KNOWN/VERIFIED evidence → capped at 7 []

VERDICT: BUILD (≥70) / TEST (50-69) / DEFER (35-49) / KILL (<35) → [verdict]
If TEST – test card: hypothesis: [] budget cap (ESTIMATE): £[] KPI: [] decision date: []
If DEFER – trigger to re-score: []
If KILL – reason in one sentence and what (if anything) re-enters differently stated: []
Language check: does pursuing this require any statement prohibited by /FACTS_BASE.md §4? [yes → KILL as stated / no]
Logged in decision log as: D-[nnn] Next re-score: [date]

Priorities / Risks / Next actions

Priorities

1. Act on the verdicts as they stand: advisory at full quota; the three TEST ideas run as the bounded experiments already dated in the 90-day plan; nothing spent on DEFER items beyond a waitlist and a note; KILL items retired from every external document.
2. Score the re-scoped Ekiti option and the grant bid/no-bid on this sheet before either goes to the board — no government proposal and no bid without a logged score.
3. Re-score all eight ideas in the week of Mon 14 December 2026 with the evidence gathered by Day 90 (interviews, audit, IP outcome, first engagements, report downloads).

Risks

Risk	Likelihood	Impact	Mitigation
Scores treated as facts (they are desk ESTIMATES by one scorer)	Medium	Medium	Evidence labels mandatory; a second scorer (adviser or NED) from the first advisory-board meeting; changes ≥ 3 points must cite new evidence
Sunk-cost pressure to keep the Ekiti proposal or the AI score alive as stated	Medium	High	Knock-out rules are written and board-visible; the re-entry routes (G-D option; v3 research track) give the ideas a legitimate home
TEST verdicts drifting into open-ended work without decision dates	High	Medium	Every TEST has a test card with a date; the Friday review lists them; the quarterly re-score closes them
Founder leverage scored optimistically because the founder is the scorer	Medium	Medium	Time audit data from the founder OS replaces the assumption after four weeks
Unweighted criteria hide that strategic fit and regulatory risk matter more at this stage	Low	Medium	Knock-outs K1 and K2 carry that weight explicitly; weights are revisited at the first quarterly re-score if the portfolio changes

Next actions

Action	Owner	Deadline
Ratify the eight verdicts and the framework at the first board meeting; log as D-007	Founder; board	Tue 13 Oct 2026
Score the grant bid/no-bid candidates from the grant calendar	Founder	Fri 23 Oct 2026
Score the re-scoped, partner-aware Ekiti option (if the ministry identifies a gap) ahead of gate G-D	Founder; advisory board profile 4 if appointed	Thu 12 Nov 2026
Score the MVP build / build-lite / wait options with the audit, IP and design-partner evidence	Founder; board	Tue 8 Dec 2026
Quarterly re-score of all live ideas; publish the scoring template as the Week 12 practical post	Founder	Week of Mon 14 Dec 2026